

Data manipulations and linear regression

Daniel J. Eck

Background

This lecture is meant to supplement Chapter 2 in your textbook.
We present a brief overview of linear regression.

The dplyr package within tidyverse

```
#install.packages("tidyverse")  
library(tidyverse)
```

dplyr provides comprehensive tools for data manipulations (or wrangling). The five main “verbs” are:

- ▶ `select()`: choose from a subset of the columns
- ▶ `filter()`: choose a subset of the rows based on logical criteria
- ▶ `arrange()`: sort the rows based on values of the columns.
- ▶ `mutate()`: add or modify the definitions of the column, and create columns that are functions of existing columns.
- ▶ `summarize()`: collapse a data frame down to a single row (per group) by aggregating vectors into a single value. Often used in conjunction with `group_by()`

The pipe operator

The pipe operator `%>%` allows for verbs to be strung in succession so that complicated manipulations can be combined within a single easily digestible sentence.

```
data %>%  
  inner_function() %>%  
  outer_function()
```

Example: Runs differential regression

```
library(Lahman)
data(Teams)
head(Teams, 3)
```

```
##   yearID lgID teamID franchID divID Rank  G  Ghome  W  L DivWin WCWin LgWin
## 1  1871   NA   BS1      BNA  <NA>    3 31    NA 20 10  <NA> <NA>    N
## 2  1871   NA   CH1      CNA  <NA>    2 28    NA 19  9  <NA> <NA>    N
## 3  1871   NA   CL1      CFC  <NA>    8 29    NA 10 19  <NA> <NA>    N
##   WSWin  R  AB  H X2B X3B HR BB SO SB CS HBP SF  RA  ER  ERA  CG SHO SV
## 1 <NA> 401 1372 426 70 37 3 60 19 73 16  NA NA 303 109 3.55 22  1  3
## 2 <NA> 302 1196 323 52 21 10 60 22 69 21  NA NA 241 77 2.76 25  0  1
## 3 <NA> 249 1186 328 35 40 7 26 25 18  8  NA NA 341 116 4.11 23  0  0
##   IPouts  HA HRA BBA SOA  E DP  FP  name
## 1  828 367  2  42  23 243 24 0.834  Boston Red Stockings
## 2  753 308  6  28  22 229 16 0.829  Chicago White Stockings
## 3  762 346 13  53  34 234 15 0.818  Cleveland Forest Citys
##   park attendance BPF PPF teamIDBR teamIDlahman45
## 1      South End Grounds I      NA 103 98      BOS      BS1
## 2      Union Base-Ball Grounds      NA 104 102      CHI      CH1
## 3 National Association Grounds      NA 96 100      CLE      CL1
##   teamIDretro
## 1      BS1
## 2      CH1
## 3      CL1
```



```

dat <- Teams %>%
  select(yearID, franchID, W, L, AB, H, X2B, X3B, HR, BB, HBP, SF,
         HA, HRA, BBA, SOA, IPouts, FP, R, RA) %>%
  filter(yearID >= 1900) %>%
  replace_na(list(HBP = 0, SF = 0)) %>%
  mutate(RD = (R - RA) / (W + L), X1B = H - (X2B + X3B + HR)) %>%
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) %>%
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) %>%
  mutate(OPS = OBP + SLG) %>%
  mutate(WHIP = 3*(HA + BBA) / IPouts) %>%
  mutate(FIP = 3*(13*HRA + 3*BBA - 2*SOA) / IPouts)
head(dat, 3)

```

```

##   yearID franchID W  L   AB    H X2B X3B HR  BB HBP SF   HA HRA BBA SOA IPouts
## 1   1900      LAD 82 54 4860 1423 199   81 26 421  81  0 1370  30 405 300  3677
## 2   1900      ATL 66 72 4952 1403 163   68 48 395  45  0 1263  59 463 340  3721
## 3   1900      CHC 65 75 4907 1276 202   51 33 343  65  0 1375  21 324 357  3813
##           FP   R  RA           RD   X1B           OBP           SLG           OPS           WHIP
## 1 0.948 816 722  0.6911765 1117 0.3590078 0.3831276 0.7421354 1.448191
## 2 0.953 778 739  0.2826087 1124 0.3418027 0.3727787 0.7145813 1.391561
## 3 0.933 635 751 -0.8285714  990 0.3168391 0.3421643 0.6590034 1.336743
##           FIP
## 1 0.8199619
## 2 1.1900027
## 3 0.4177813

```


Note: other packages may contain functions with the same name as those in `dplyr`. For example, the `MASS` package also contains a `select` function.

In the event that you have both `dplyr` and `MASS` loaded in an R session, you can access `dplyr`'s `select` function using `dplyr::select`

```
dat <- Teams %>%
  dplyr::select(yearID, franchID, W, L, AB, H, X2B, X3B, HR, BB, HBP, SF,
               HA, HRA, BBA, SOA, IPouts, FP, R, RA) %>%
  filter(yearID >= 1900) %>%
  replace_na(list(HBP = 0, SF = 0)) %>%
  mutate(RD = (R - RA) / (W + L), X1B = H - (X2B + X3B + HR)) %>%
  mutate(OBP = (H + BB + HBP) / (AB + BB + HBP + SF)) %>%
  mutate(SLG = (X1B + 2*X2B + 3*X3B + 4*HR) / AB) %>%
  mutate(OPS = OBP + SLG) %>%
  mutate(WHIP = 3*(HA + BBA) / IPouts) %>%
  mutate(FIP = 3*(13*HRA + 3*BBA - 2*SOA) / IPouts)
```

Baseball is a game of offense, pitching, and defense. Let's see how well runs differential per game is explained by:

- ▶ **OPS**: on base percentage plus slugging percentage
- ▶ **WHIP**: walks and hits allowed divided by innings pitched
- ▶ **FP**: fielding percentage

using a linear regression model

```
m <- lm(RD ~ OPS + WHIP + FP, data = dat)
```

Regression review

Regression model:

$$y = \beta_1 x_1 + \cdots + \beta_p x_p + \varepsilon; \quad \varepsilon \sim N(0, \sigma^2),$$

where we usually specify a model intercept by setting $x_1 = 1$.

Can also write in vector notation:

$$y = \mathbf{x}'\beta + \varepsilon; \quad \varepsilon \sim N(0, \sigma^2),$$

where $\mathbf{x}, \beta \in \mathbb{R}^p$.

Either way, this model relies on a few assumptions:

- ▶ a linear relationship is present
- ▶ errors are independent and identically distributed
- ▶ errors are normally distributed mean 0 and common variance σ^2

Regression review

Remember that linear regression is about modeling a conditional expectation, the scattering of points is noise. Interest is in

$$E(y|\mathbf{x}) = \mathbf{x}'\beta,$$

where it is important to choose the variables comprising \mathbf{x} and to be able to defend those choices.

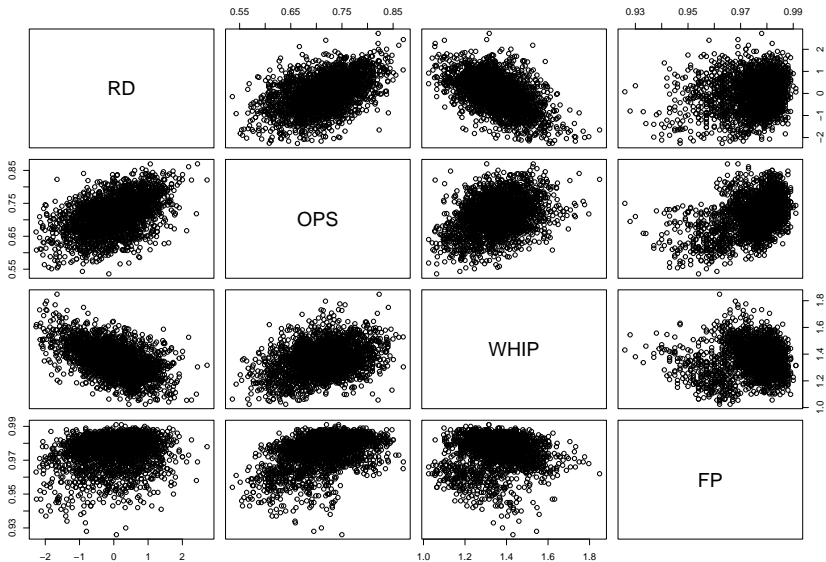
Yes, baseball IS a game of offense, pitching, and defense.

```
summary(m)
```

```
##
## Call:
## lm(formula = RD ~ OPS + WHIP + FP, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.46029 -0.21779 -0.00699  0.20934  1.48506
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  13.41760    0.79402   16.90  <2e-16 ***
## OPS          11.81943    0.15773   74.93  <2e-16 ***
## WHIP         -5.44368    0.06515  -83.55  <2e-16 ***
## FP          -14.84942    0.84580  -17.56  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3513 on 2606 degrees of freedom
## Multiple R-squared:  0.787, Adjusted R-squared:  0.7868
## F-statistic: 3210 on 3 and 2606 DF, p-value: < 2.2e-16
```





Linearity holds.

```
pairs(dat %>% select(RD, OPS, WHIP, FP))
```



Normality of mean-zero errors with constant variance holds.
Although a slightly heavy right tail is observed in the residuals.

```
par(mfrow = c(2,2))  
plot(m)
```



The compiled fractions should be roughly 68% and 95% if errors are truly normal with common variance.

```
library(broom)
dat_aug <- augment(m, data = dat)
dat_aug %>%
  mutate(rmse = sqrt((mean(.resid^2)))) %>%
  summarise(N = n(),
            within_1rmse = sum(abs(.resid) < rmse),
            within_2rmse = sum(abs(.resid) < 2 * rmse)) %>%
  mutate(within_1rmse_pct = within_1rmse / N,
         within_2rmse_pct = within_2rmse / N)

## # A tibble: 1 x 5
##       N within_1rmse within_2rmse within_1rmse_pct within_2rmse_pct
##   <int>      <int>      <int>      <dbl>      <dbl>
## 1  2610      1844      2469      0.707      0.946
```


We will suppose that independence holds, or that any violations of this assumption that may be present in this data do not materially effect our overall conclusions.

A saturated model (one parameter per observation) does not fit the data better than our model with three variables and an intercept.

```
# likelihood ratio test of fitted model vs a saturated model  
m_glm <- glm(RD ~ OPS + WHIP + FP, data = dat)  
pchisq(m_glm$deviance, m_glm$df.residual, lower = FALSE)
```

```
## [1] 1
```

Thus we have a well-fitting simple and useful model that provides satisfactory dimension reduction.

Investigate large residuals

```
dat_aug %>% filter(abs(.resid) >= 1) %>%  
  select(yearID, franchID, W, RD, OPS, WHIP, FP, .resid, .fitted) %>%  
  mutate(across(4:9, round, 3)) %>%  
  arrange(desc(.resid))
```

```
## # A tibble: 27 x 9  
##   yearID franchID      W      RD      OPS      WHIP      FP .resid .fitted  
##   <int> <fct>    <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1  1949 NYY         97 1.25  0.758  1.49 0.977  1.48 -0.238  
## 2  1936 NYY        102 2.18  0.861  1.53 0.973  1.34  0.838  
## 3  1939 NYY        106 2.72  0.821  1.32 0.978  1.29  1.43  
## 4  1950 BOS         94 1.45  0.846  1.59 0.981  1.24  0.211  
## 5  1949 BOS         96 1.49  0.8    1.48 0.98   1.22  0.27  
## 6  1935 DET         93 1.68  0.798  1.44 0.979  1.20  0.487  
## 7  1948 BOS         96 1.21  0.779  1.48 0.981  1.19  0.021  
## 8  1950 NYY         98 1.45  0.804  1.48 0.979  1.12  0.329  
## 9  1943 CIN         87 0.422 0.653  1.34 0.98   1.12 -0.694  
## 10 1914 OAK         99 1.45  0.693  1.27 0.966  1.10  0.349  
## # ... with 17 more rows
```


Investigate large fitted values

```
dat_aug %>% filter(.fitted >= 2) %>%  
  select(yearID, franchID, W, RD, OPS, WHIP, FP, .resid, .fitted)
```

```
## # A tibble: 4 x 9  
##   yearID franchID      W      RD      OPS  WHIP      FP .resid .fitted  
##   <int> <fct>    <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1  1927 NYY      110  2.44  0.870  1.30  0.969  0.231  2.21  
## 2  2019 HOU      107  1.73  0.848  1.13  0.988 -0.886  2.61  
## 3  2019 LAD      106  1.69  0.810  1.10  0.982 -0.722  2.41  
## 4  2020 LAD       43  2.27  0.821  1.06  0.982 -0.518  2.79
```


A closer look at problems with fit

```
qqnorm(resid(m)); qqline(resid(m))  
abline(a=0.5, b=0, lty = 2, col = "red")
```



A closer look at problems with fit

```
plot(table(dat_aug %>% filter(abs(.resid) >= 0.5) %>%  
  pull(yearID)), ylab = "number of model under counts")
```



League conditions change over time

