

Value of plays using run expectancy

Daniel J. Eck

Background

This lecture is meant to supplement Chapter 5 in your textbook.
We will now study the value of baseball events.

Introduction to the run expectancy matrix

The run expectancy matrix is the average number of runs scored for each combination of outs and runners on base.

There are 8 possible arrangements of runners on the three bases, and the number of possible outs equals 3. Thus there are a total of 24 possible arrangements of outs and runners which form the run expectancy matrix.

The run expectancy matrix establishes a baseline value for baseball events in a context-free environment.

Data

We will calculate the run expectancy matrix. We first load in relevant software packages

```
library(Lahman)
library(tidyverse)
library(retrosheet)
```

and then we load in play-by-play data from the 2016 season

```
fields <- read_csv("fields.csv")
dat2016 <- read_csv("all2016.csv",
                    col_names = pull(fields, Header),
                    na = character())
```

Data manipulations

We compute a runs scored in the remainder of the inning variable and add it to `dat2016`.

```
dat2016 <- dat2016 %>%
  mutate(RUNS = AWAY_SCORE_CT + HOME_SCORE_CT,
         HALF.INNING = paste(GAME_ID, INN_CT, BAT_HOME_ID),
         RUNS.SCORED = (BAT_DEST_ID > 3) + (RUN1_DEST_ID > 3) +
           (RUN2_DEST_ID > 3) + (RUN3_DEST_ID > 3))

half_innings <- dat2016 %>%
  group_by(HALF.INNING) %>%
  summarise(Outs.Inning = sum(EVENT_OUTS_CT),
           Runs.Inning = sum(RUNS.SCORED),
           Runs.Start = first(RUNS),
           MAX.RUNS = Runs.Inning + Runs.Start)

dat2016 <- dat2016 %>%
  inner_join(half_innings, by = "HALF.INNING") %>%
  mutate(RUNS.ROI = MAX.RUNS - RUNS)
```

Creating the matrix

Now that the runs scored in the remainder of the inning variable have been computed for each plate appearance, it is straightforward to compute the run expectancy matrix.

We create a `BASES` variable which indicates the base runner state (eg, 100 corresponds to a runner on first), and a `STATE` variable which adds the number of outs to `BASES`.

```
dat2016 <-  
  dat2016 %>% mutate(BASES = paste(ifelse(BASE1_RUN_ID != "", 1, 0),  
                                   ifelse(BASE2_RUN_ID != "", 1, 0),  
                                   ifelse(BASE3_RUN_ID != "", 1, 0), sep = ""),  
                     STATE = paste(BASES, OUTS_CT))
```

Creating the matrix (continued)

We now trim `dat2016` to only include plays in which the state of the game changed and a half inning reached 3 outs.

```
dat2016 <- dat2016 %>%
  mutate(NRUNNER1 = as.numeric(RUN1_DEST_ID == 1 | BAT_DEST_ID == 1),
         NRUNNER2 = as.numeric(RUN1_DEST_ID == 2 | RUN2_DEST_ID == 2 | BAT_DEST_ID == 2),
         NRUNNER3 = as.numeric(RUN1_DEST_ID == 3 | RUN2_DEST_ID == 3 |
                                RUN3_DEST_ID == 3 | BAT_DEST_ID == 3),
         NOUTS = OUTS_CT + EVENT_OUTS_CT,
         NEW.BASES = paste(NRUNNER1, NRUNNER2, NRUNNER3, sep = ""),
         NEW.STATE = paste(NEW.BASES, NOUTS)) %>%
  filter((STATE != NEW.STATE) | (RUNS.SCORED > 0)) %>%
  filter(Outs.Inning == 3)
```

Creating the matrix (continued)

We now create the run expectancy matrix RUNS_out

```
RUNS <- dat2016 %>%
  group_by(STATE) %>%
  summarize(Mean = mean(RUNS.ROI)) %>%
  mutate(Outs = substr(STATE, 5, 5)) %>%
  arrange(Outs)

RUNS_out <- matrix(round(RUNS$Mean, 2), 8, 3)
dimnames(RUNS_out)[[1]] <- c("000", "001", "010", "011",
                             "100", "101", "110", "111")
dimnames(RUNS_out)[[2]] <- c("0 outs", "1 out", "2 outs")
RUNS_out
```

```
##      0 outs 1 out 2 outs
## 000    0.50 0.27 0.11
## 001    1.35 0.94 0.37
## 010    1.13 0.67 0.31
## 011    1.93 1.36 0.55
## 100    0.86 0.51 0.22
## 101    1.72 1.20 0.48
## 110    1.44 0.92 0.41
## 111    2.11 1.54 0.70
```


Measuring the success of a batting play

When a player comes to bat with a particular runners out situation, the run expectancy matrix tells us the number of runs a team is expected to score in the remainder of the half inning:

$$\text{Run Value} = (\text{Runs}_{\text{new state}} - \text{Runs}_{\text{old state}}) + \text{Runs}_{\text{scored on play}}$$

```
dat2016 <- dat2016 %>%  
  left_join(select(RUNS, - Outs), by = "STATE") %>%  
  rename(Runs.State = Mean) %>%  
  left_join(select(RUNS, -Outs), by = c("NEW.STATE" = "STATE")) %>%  
  rename(Runs.New.State = Mean) %>%  
  replace_na(list(Runs.New.State = 0)) %>%  
  mutate(run_value = Runs.New.State - Runs.State + RUNS.SCORED)
```

Example: Jose Altuve

We will now study Jose Altuve's 2016 season.

The code below isolates the run value for each of Altuve's batting events and displays his first 3 batting events.

```
data('People')
altuve.id <- People %>% filter(nameFirst == "Jose", nameLast == "Altuve") %>% pull(retroID)
# BAT_EVENT_FL == TRUE distinguishes batting events from non batting events like steals
altuve <- dat2016 %>% filter(BAT_ID == altuve.id,
                           BAT_EVENT_FL == TRUE)
altuve %>% select(STATE, NEW.STATE, run_value) %>%
  slice(1:3)
```

```
## # A tibble: 3 x 3
##   STATE NEW.STATE run_value
##   <chr> <chr>      <dbl>
## 1 000 1 000 2      -0.162
## 2 000 1 100 1       0.244
## 3 000 1 000 2      -0.162
```

We can see that Jose Altuve was 13th in total RE24 value.

```
dat2016 %>% inner_join(People %>% select(nameFirst, nameLast, retroID),  
  by = c("BAT_ID" = "retroID")) %>%  
  filter( BAT_EVENT_FL == TRUE) %>%  
  group_by(BAT_ID) %>%  
  summarise(nameFirst = unique(nameFirst),  
    nameLast = unique(nameLast),  
    RE24 = sum(run_value())) %>%  
  arrange(desc(RE24)) %>% as.data.frame() %>% head(20)
```

##	BAT_ID	nameFirst	nameLast	RE24
## 1	troum001	Mike	Trout	65.21086
## 2	ortid001	David	Ortiz	58.75680
## 3	freef001	Freddie	Freeman	46.19011
## 4	donaj001	Josh	Donaldson	46.10413
## 5	vottj001	Joey	Votto	45.74782
## 6	bryak001	Kris	Bryant	45.68620
## 7	arenn001	Nolan	Arenado	45.34225
## 8	murpd006	Daniel	Murphy	44.57236
## 9	rizza001	Anthony	Rizzo	40.19773
## 10	bettm001	Mookie	Betts	35.65014
## 11	cabrm001	Miguel	Cabrera	34.27806
## 12	goldp001	Paul	Goldschmidt	34.27474
## 13	altuj001	Jose	Altuve	33.16617
## 14	enca001	Edwin	Encarnacion	32.03130
## 15	blacc001	Charlie	Blackmon	31.86506
## 16	ramih003	Hanley	Ramirez	31.39952
## 17	machm001	Manny	Machado	29.48961
## 18	belta001	Adrian	Beltre	29.11069
## 19	gonzc001	Carlos	Gonzalez	28.29142
## 20	cruzn002	Nelson	Cruz	28.18916

We can see the number of opportunities Jose Altuve had in each base out state.

```
altuve %>% group_by(STATE) %>%  
  summarise(N = n(), avg_run_value = mean(run_value),  
            total_run_value = sum(run_value), se_run_value = sd(run_value)/ sqrt(N))  
as.data.frame()
```

##	STATE	N	avg_run_value	total_run_value	se_run_value
## 1	000 0	185	0.036051610	6.6695479	0.02707034
## 2	000 1	104	0.030673539	3.1900481	0.02786826
## 3	000 2	126	-0.002899510	-0.3653383	0.01736461
## 4	001 0	3	-0.103059751	-0.3091793	0.30703088
## 5	001 1	8	0.382615911	3.0609273	0.15351209
## 6	001 2	12	0.100203614	1.2024434	0.14924366
## 7	010 0	15	-0.078676867	-1.1801530	0.15080076
## 8	010 1	22	0.016366766	0.3600688	0.13704177
## 9	010 2	22	0.049357251	1.0858595	0.10129068
## 10	011 0	5	0.317988565	1.5899428	0.37330473
## 11	011 1	5	-0.096378018	-0.4818901	0.18380432
## 12	011 2	8	0.290682347	2.3254588	0.34111391
## 13	100 0	31	0.162685083	5.0432376	0.13171672
## 14	100 1	58	0.085061807	4.9335848	0.07587254
## 15	100 2	39	0.005568296	0.2171636	0.06065280
## 16	101 0	8	-0.316187264	-2.5294981	0.20184533
## 17	101 1	7	0.187174415	1.3102209	0.21187277
## 18	101 2	6	0.363020367	2.1781222	0.23086490
## 19	110 0	6	0.684712857	4.1082771	0.50690064
## 20	110 1	14	-0.060673082	-0.8494231	0.22048071
## 21	110 2	19	0.089659990	1.7035398	0.19236237
## 22	111 0	2	-0.239618713	-0.4792374	0.32927105
## 23	111 1	5	0.215543377	1.0777169	0.52456637
## 24	111 2	1	-0.695272354	-0.6952724	NA

Cleaner presentation of sample sizes for each out base state:

```
altuve_RE <- altuve %>% group_by(STATE) %>%  
  summarize(N = n(), avg_run_value = mean(run_value)) %>%  
  mutate(Outs = substr(STATE, 5, 5)) %>%  
  arrange(Outs)  
  
altuve_N_mat <- matrix(round(altuve_RE$N, 4), 8, 3)  
dimnames(altuve_N_mat)[[1]] <- c("000", "001", "010", "011",  
                                   "100", "101", "110", "111")  
dimnames(altuve_N_mat)[[2]] <- c("0 outs", "1 out", "2 outs")  
altuve_N_mat
```

```
##      0 outs 1 out 2 outs  
## 000    185   104   126  
## 001      3     8    12  
## 010     15    22    22  
## 011      5     5     8  
## 100     31    58    39  
## 101      8     7     6  
## 110      6    14    19  
## 111      2     5     1
```

```
colSums(altuve_N_mat)
```

```
## 0 outs 1 out 2 outs  
##    255    223    233
```

```
rowSums(altuve_N_mat)
```

```
## 000 001 010 011 100 101 110 111  
## 415  23  59  18 128  21  39   8
```

Cleaner presentation of run value for each out base state:

```
altuve_RE_mat <- matrix(round(altuve_RE$avg_run_value, 4), 8, 3)
dimnames(altuve_RE_mat)[[1]] <- c("000", "001", "010", "011",
                                   "100", "101", "110", "111")
dimnames(altuve_RE_mat)[[2]] <- c("0 outs", "1 out", "2 outs")

colMeans(altuve_RE_mat)
```

```
##      0 outs      1 out      2 outs
## 0.0579875 0.0950500 0.0250500
rowMeans(altuve_RE_mat)
```

```
##           000           001           010           011           100           101
## 0.02130000  0.12656667 -0.00430000  0.17076667  0.08446667  0.07800000
##           110           111
## 0.23790000 -0.23980000
```

We detect statistically significant differences for Jose Altuve's performance across base out states using an anova test. However, a close look reveals that these differences are not intuitive and our detected statistical significance may be just noise.

```
# performance in different states
summary(aov(run_value ~ -1 + STATE, data = altuve))
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## STATE          24    9.97   0.4153    1.873 0.00716 **
## Residuals 687 152.32   0.2217
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
round(coef(summary(lm(run_value ~ -1 + STATE, data = altuve))), 3)
```

##		Estimate	Std. Error	t value	Pr(> t)
##	STATE000 0	0.036	0.035	1.041	0.298
##	STATE000 1	0.031	0.046	0.664	0.507
##	STATE000 2	-0.003	0.042	-0.069	0.945
##	STATE001 0	-0.103	0.272	-0.379	0.705
##	STATE001 1	0.383	0.166	2.298	0.022
##	STATE001 2	0.100	0.136	0.737	0.461
##	STATE010 0	-0.079	0.122	-0.647	0.518
##	STATE010 1	0.016	0.100	0.163	0.871
##	STATE010 2	0.049	0.100	0.492	0.623
##	STATE011 0	0.318	0.211	1.510	0.131
##	STATE011 1	-0.096	0.211	-0.458	0.647
##	STATE011 2	0.291	0.166	1.746	0.081
##	STATE100 0	0.163	0.085	1.924	0.055
##	STATE100 1	0.085	0.062	1.376	0.169
##	STATE100 2	0.006	0.075	0.074	0.941
##	STATE101 0	-0.316	0.166	-1.899	0.058
##	STATE101 1	0.187	0.178	1.052	0.293
##	STATE101 2	0.363	0.192	1.888	0.059
##	STATE110 0	0.685	0.192	3.562	0.000
##	STATE110 1	-0.061	0.126	-0.482	0.630
##	STATE110 2	0.090	0.108	0.830	0.407
##	STATE111 0	-0.240	0.333	-0.720	0.472
##	STATE111 1	0.216	0.211	1.024	0.306
##	STATE111 2	-0.695	0.471	-1.477	0.140

Two-way anova tests do not reveal statistical significance.

```
summary(aov(run_value ~ -1 + BASES + Outs, data = altuve %>%  
  mutate(Outs = substr(STATE, 5, 5))))
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## BASES         8   3.04   0.3802   1.675  0.101  
## Outs          2   0.15   0.0774   0.341  0.711  
## Residuals    701 159.10   0.2270
```

```
summary(aov(run_value ~ -1 + BASES + Outs, data = altuve %>%  
  filter(!(BASES %in% c("111", "011")))) %>%  
  mutate(Outs = substr(STATE, 5, 5))))
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)  
## BASES         6   2.39   0.3975   1.896 0.0792 .  
## Outs          2   0.16   0.0816   0.389 0.6777  
## Residuals    677 141.99   0.2097  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

A dichotomy between RISP and no RISP reveals statistical significance.

```
# performance with RISP
altuve <- altuve %>% mutate(RISP = ifelse(!BASES %in% c("100", "000"), 1, 0))
summary(aov(run_value ~ -1 + RISP + Outs, data = altuve %>%
  mutate(Outs = substr(STATE, 5, 5))))

##              Df Sum Sq Mean Sq F value Pr(>F)
## RISP          1   1.08  1.0813    4.766 0.0294 *
## Outs          3   0.81  0.2687    1.184 0.3147
## Residuals    707 160.41  0.2269
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
altuve %>% group_by(RISP) %>% summarise(N = n(), avg_run_value = mean(run_value),
  total_run_value = sum(run_value), se_run_value = sd(run_value) / sqrt(N))

## # A tibble: 2 x 5
##   RISP      N avg_run_value total_run_value se_run_value
##   <dbl> <int>         <dbl>         <dbl>         <dbl>
## 1     0   543         0.0363          19.7         0.0164
## 2     1   168         0.0802          13.5         0.0537
```


However, Jose Altuve ranked 75th among roughly 200 full time players in the difference in run value with RISP and with runners not in RISP.

We have mixed results. Can check out [brief for a complete breakdown](#)

```
dat2016 %>% inner_join(People %>% select(nameFirst, nameLast, retroID),
  by = c("BAT_ID" = "retroID")) %>%
  filter( BAT_EVENT_FL == TRUE) %>%
  mutate(RISP = ifelse(!BASES %in% c("100", "000"), 1, 0)) %>%
  group_by(BAT_ID) %>%
  summarise(nameFirst = unique(nameFirst),
    nameLast = unique(nameLast),
    N = n(),
    diff = mean(run_value[which(RISP == 1)]) - mean(run_value[which(RISP == 0)])) %>%
  filter(N >= 400) %>% arrange(desc(diff)) %>% slice(73:77)
```

```
## # A tibble: 5 x 5
##   BAT_ID  nameFirst nameLast      N    diff
##   <chr>   <chr>    <chr>   <int> <dbl>
## 1 martv001 Victor    Martinez  608 0.0454
## 2 keplm001 Max      Kepler    445 0.0448
## 3 altuj001 Jose     Altuve    711 0.0440
## 4 lemad001 DJ       LeMahieu  634 0.0427
## 5 ortid001 David    Ortiz     623 0.0415
```

```
ggplot(altuve, aes(BASES, run_value)) +
  geom_jitter(width = 0.15, alpha = 0.20) +
  geom_hline(yintercept = 0, color = "blue") +
  xlab("Runners") + ylab("Run value")
```



Example: All batters

We create a new variable: total starting runs potential
Runs.Start.


```
runs <- dat2016 %>%
  filter(BAT_EVENT_FL == TRUE) %>%
  inner_join(People, by = c("BAT_ID" = "retroID")) %>%
  group_by(BAT_ID) %>%
  summarise(RE24 = sum(run_value),
            PA = length(run_value),
            Runs.Start = sum(Runs.State),
            nameLast = unique(nameLast)) %>%
  filter(PA >= 400)

head(runs)
```

```
## # A tibble: 6 x 5
##   BAT_ID    RE24    PA Runs.Start nameLast
##   <chr>    <dbl> <int>      <dbl> <chr>
## 1 abrej003  12.3   693      334. Abreu
## 2 alony001  -6.94   528      247. Alonso
## 3 altuj001  33.2   711      341. Altuve
## 4 andet001 -10.9   428      202. Anderson
## 5 andre001  17.4   567      256. Andrus
## 6 aokin001  -2.12  466      229. Aoki
```

Batters with larger values of `Runs.Start` tend to have larger runs contributions. Batters with at least 40 RE24 are labeled.

```
library(ggrepel)
ggplot(runs, aes(Runs.Start, RE24)) +
  geom_point() + geom_smooth() +
  geom_hline(yintercept = 0, color = "blue") +
  geom_text_repel(data = filter(runs, RE24 >= 40), aes(label = nameLast))
```

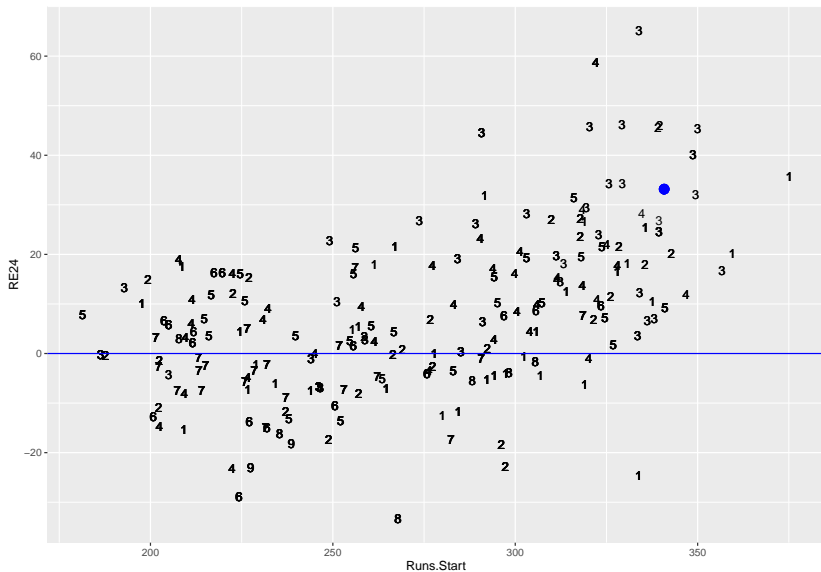


Simple lineup analysis

Managers like to put their best hitters near the middle of the lineup.

```
regulars <- dat2016 %>% inner_join(runs, by = "BAT_ID")
positions <- regulars %>% group_by(BAT_ID, BAT_LINEUP_ID) %>%
  summarise(N = n()) %>% arrange(desc(N)) %>%
  mutate(Position = first(BAT_LINEUP_ID))
runs <- runs %>% inner_join(positions, by = "BAT_ID")

ggplot(runs, aes(Runs.Start, RE24, label = Position)) +
  geom_text() +
  geom_hline(yintercept = 0, color = "blue") +
  geom_point(data = filter(runs, BAT_ID == altuve.id),
    size = 4, shape = 16, color = "blue")
```



Value of home runs

```
## get home runs
home_runs <- dat2016 %>% filter(EVENT_CD == 23)

home_runs_N <- home_runs %>% group_by(STATE) %>%
  mutate(Outs = substr(STATE, 5, 5)) %>%
  arrange(Outs) %>%
  summarise(Outs = unique(Outs), N = n()) %>%
  arrange(Outs)

home_runs_N_mat <- matrix(round(home_runs_N$N / sum(home_runs_N$N), 3), 8, 3)
dimnames(home_runs_N_mat)[[1]] <- c("000", "001", "010", "011",
                                     "100", "101", "110", "111")
dimnames(home_runs_N_mat)[[2]] <- c("0 outs", "1 out", "2 outs")
home_runs_N_mat
```


```
##      0 outs 1 out 2 outs
## 000  0.272 0.172 0.150
## 001  0.002 0.007 0.011
## 010  0.018 0.027 0.028
## 011  0.004 0.007 0.007
## 100  0.057 0.064 0.061
## 101  0.005 0.013 0.011
## 110  0.015 0.023 0.028
## 111  0.003 0.008 0.008
```

```
avg_hr <- home_runs %>% summarise(avg_run_value = mean(run_value))
avg_hr
```

```
## # A tibble: 1 x 1
##   avg_run_value
##   <dbl>
## 1         1.38
```

```
ggplot(home_runs, aes(run_value)) +
  geom_histogram() +
  geom_vline(data = avg_hr, aes(xintercept = avg_run_value),
            color = "blue", size = 1.5) +
  annotate("text", 1.6, 2000, label = "Mean Run\nValue", color = "blue")
```

```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
```



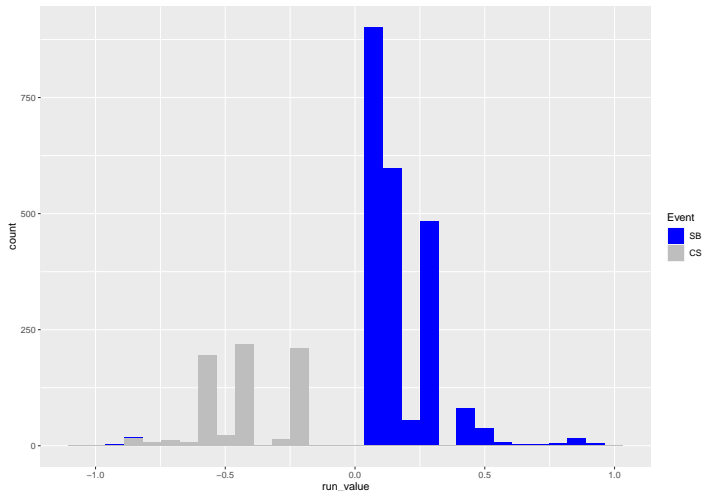
Value of base stealing

```
stealing <- dat2016 %>% filter(EVENT_CD %in% c(4,6))
stealing %>% group_by(EVENT_CD) %>% summarise(N = n(),
  avg_run_value = mean(run_value)) %>%
  mutate(pct = N/sum(N))
```

```
## # A tibble: 2 x 4
##   EVENT_CD      N avg_run_value  pct
##   <dbl> <int>      <dbl> <dbl>
## 1       4  2199         0.180 0.756
## 2       6   710        -0.416 0.244
```

Histogram of the run values of all steal attempts during the 2016 season.

```
ggplot(stealing, aes(run_value, fill = factor(EVENT_CD))) +  
  geom_histogram() +  
  scale_fill_manual(name = "Event", values = c("blue", "grey"),  
                    labels = c("SB", "CS"))
```



We can compute the marginal break-even success rate needed to justify a stolen base attempt across the 2016 season

$$a * SB_{\text{avg value}} + (1 - a) * CS_{\text{avg value}} = 0$$

which implies that

$$a = - \frac{CS_{\text{avg value}}}{SB_{\text{avg value}} - CS_{\text{avg value}}}.$$

From a previous slide we compute

$$a = 0.416 / (0.180 + 0.416)$$

```
## [1] 0.6979866
```